RESEARCH POSTER

What Does an LLM Judge Actually Measure?

The Role of Evaluation Frames in Construct Discrimination

Brian Kerrigan | Independent Researcher

REVIEW DRAFT v0.2

Target-bound scoring established target-defect discrimination for all three judges.

The same prerequisite was not established under ordinary holistic scoring in this fixed corpus.

01 / THE REAL-WORLD BUG

A correct alert, a wrong reading

A solar-alert rule fired correctly. Its output omitted the operative threshold; the downstream agent used a stale label and de-escalated the warning.

A 48-call exploratory replay raised a question: **does a judge distinguish a defect when another is present?**

The incident and replay contributed no S1 or S2 observations.

02 / THE CONTROLLED TEST

A target defect, four states

B = the target defect. A = another defect. Both are defined separately for each specimen.

00 Neither defect	**A0** A only
0B B only	**AB** Both defects

S1: 4,800 judgments
10 frozen specimens; 6 primary semantic specimens and 4 separate mechanical controls.

No judge established standalone discrimination. The confirmatory attenuation test stayed unopened.

Three fixed provider/model/settings bundles; differences are configuration-specific, not vendor effects.

03 / CHANGE THE EVALUATION FRAME

S2: same artifacts, three ways to score

8,640 fresh judgments | 6 specimens x 4 states x 3 frames x 3 judges x 40

F0 Ordinary holistic	**F1** Named holistic	**F2** Target bound
Target issue not identified One overall score	Target issue explicitly named One overall score	Same target description Score only that issue

Standalone target discrimination

Compare 00 vs 0B (A absent): probability that the B-absent artifact scores higher, with ties split equally. Higher means better separation; 0.50 means no ordering advantage.

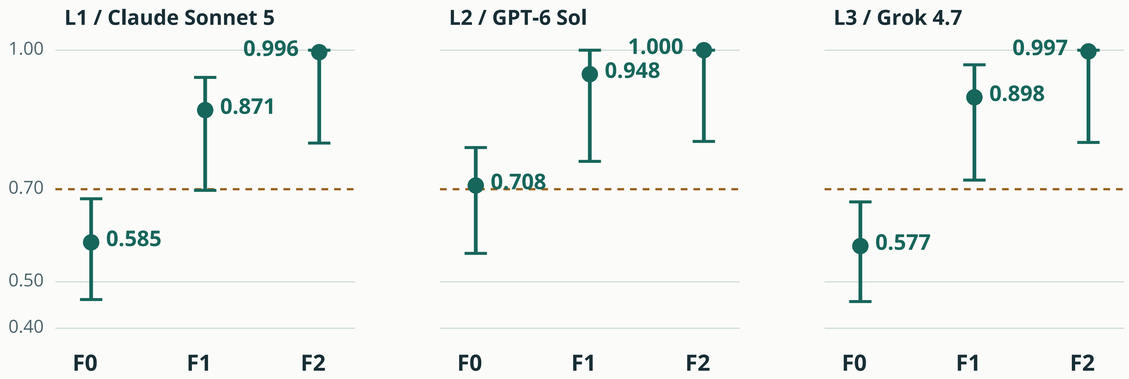

Simultaneous 90% intervals; preregistered familywise α = 0.10.
Dashed line: registered 0.70 discrimination threshold.

F1: target named, 2 of 3 channels established

Point estimates rose for all three judges; formal naming gain was not established.

The claim boundary matters. F2 success was an expected manipulation check. Neither naming gain nor incremental binding gain was established. Complete binding recovery was not supported; conditional attenuation remained unidentified.

04 / A SECOND MEASUREMENT PROBLEM

Humans disagreed about written recognition

28.3%

initial agreement: 34 / 120 rationales

Three blinded coders; initial pairs disagreed on 71.7% of items. Krippendorff's alpha = -0.143.

Poor agreement concerns the reliability of coding explicit recognition. It does not establish what the models internally recognized.

05 / WHY THE TRAIL IS AUDITABLE

Observable calls, bounded claims

Holonograph recorded call-time identity, settings provenance, usage and cache metadata, and immutable digest-only events. Private responses stayed outside Lens events.

Controllers enforced schedules and acceptance rules. Results were sealed before interpretation; synthetic tests challenged the analysis engines.

A branch map preceded S1. The exact S2 protocol was fixed after S1 and before coder outcomes. This was a sequentially prospective program.

Holonograph is developed by the author. Human coders volunteered and were employees of the author's businesses.

Before trusting a score, demonstrate that it discriminates the construct you need.

Fixed-corpus findings. Descriptive frame progression does not establish the preregistered gain claims.
Methods, full intervals, limitations and references are in the submitted paper and public evidence deposit.

Public evidence trail

github.com/inspectorSlap/attestations

WORKING SIZE: 48 x 36 in landscape | Review draft - do not print